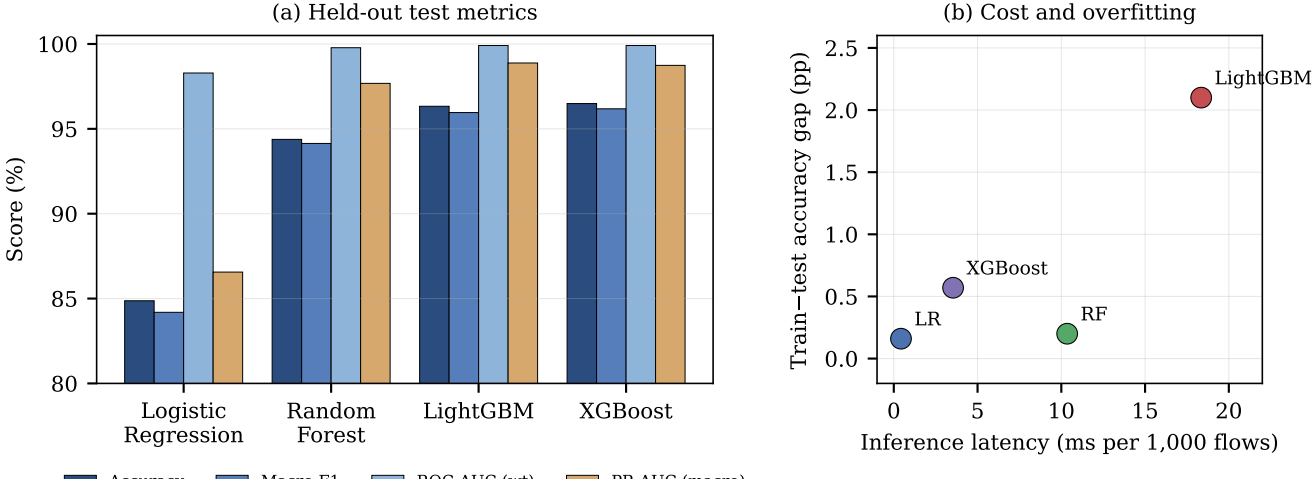
(a) Held-out test metrics
Score (%)
100
95
90
85
80
Logistic Regression
Random Forest
LightGBM
XGBoost
Accuracy
Macro-F1
ROC-AUC (wt)
PR-AUC (macro)
(b) Cost and overfitting
Train−test accuracy gap (pp)
2.5
2.0
1.5
1.0
0.5
0.0
LightGBM
XGBoost
LR
RF
0
5
10
15
20
Inference latency (ms per 1,000 flows)